

Bridge Transfers and Quarantine in a Captive Turtle Farm

CITS4403 Research Project, 2026. Cam Zhou (24349432) and Wenhao Zhang (24364285).

Code, notebook and demonstration video: <https://github.com/camzhou831-star/cits4403-research-project>
(walkthrough: `notebooks/project-walkthrough.ipynb`; *animated* *example*:
`results/demo/demo-3min.mp4`).

Abstract. A manager who can close only a few tanks of a captive facility must decide which ones. We built an agent-based model of 200 turtles in 20 tanks on a modular transfer network and ran 5,200 paired simulations, varying the transfer rate, the response delay and which two tanks are quarantined for 14 days. The transfer rate dominates: raising it from 0.01 to 0.025 multiplies the attack rate by 3.3, whereas quarantine reduces it by at most 10.6%, and targeting the highest-betweenness tanks is not consistently better than choosing at random. An exploratory analysis shows why: the targeted tanks do cover the bridges between regions (40% of between-region edges, against 21%), but outbreaks first cross regions around day 15, so a fixed 14-day closure overlaps at most half of these first crossings. When a quarantine happens matters as much as where it is placed.

1. Introduction

Moving animals between housing units is a routine part of captive husbandry, and it is also a pathway for infection. In livestock systems, recorded animal movements have been linked to the spread of bovine tuberculosis [1], and the structure of the movement network shapes how pathogens spread [2]. A facility made of many tanks has the structure of a **metapopulation**: transmission is local within each unit, and units are connected by occasional transfers. In metapopulation models, an infection reaches a large share of subpopulations only when mobility exceeds an invasion threshold [3]. More generally, the structure of a contact network shapes how an epidemic spreads [4], and in networks with strong community structure the few links between communities strongly affect whether a local outbreak becomes a global one [5].

This structure raises a practical question for anyone who can restrict movement: if only a few tanks can be closed, does it matter *which* tanks they are? In heterogeneous networks, targeting nodes by their network position is far more effective than uniform random action [6], and network measures have been used to identify individuals at high risk of infection [7]. In networks with strong community structure, interventions aimed at individuals that bridge communities outperform those aimed at highly connected individuals [5]. Betweenness centrality [8], the fraction of shortest paths passing through a node, is a natural measure of such bridging that can be computed before an outbreak.

We study these ideas in a deliberately simple system: 200 synthetic turtles in 20 tanks, grouped into 4 regions of a modular transfer network. Our primary question is how the cross-tank transfer rate and the response delay affect the final outbreak size and the number of affected tanks. Our secondary question is whether, under the same intervention budget, quarantining the highest-betweenness tanks reduces spread more than quarantining randomly chosen tanks. Before the experiment we stated three hypotheses: (H1) a low but non-zero transfer rate lets a local outbreak spread between regions, raising both outcomes; (H2) a longer response delay does not lower either outcome; and (H3) under an equal budget, highest-betweenness quarantine lowers both outcomes more than random quarantine.

Contribution. Our contribution is a controlled computational experiment, not a new simulator, and none of the unit's taught models [9] is reused. The studies above target individuals in contact networks; here the unit of control is a tank with finite capacity, so closing a tank both isolates its occupants and changes where other agents can move. We compare strategies under an identical

budget in tank-days and with common random numbers [10], so every strategy in a comparison sees exactly the same outbreak until the quarantine starts, and any later difference is caused by the choice of tanks. Finally, we use the recorded runs to explain the size of the quarantine effect (Section 4.4). The model is a stylised explanation of mechanism, not a forecast for any real facility.

2. Model

We built a discrete-time stochastic agent-based model (ABM); one time step is one day, and all data are synthetic. Each of the 200 turtle agents has a disease state, susceptible (S), infectious (I) or recovered (R), as in the SIR model [11], and lives in exactly one of 20 tanks. A tank holds at most 12 agents and is either `open` or `quarantined`. At the start, every tank holds 10 agents and one agent, chosen uniformly at random, is infected. The capacity leaves 2 free places per tank, so a transfer is possible but fails when every neighbouring tank is full. All parameters, their sources and the seeds are listed in Appendix G.

Tanks are the nodes of a static, undirected **modular transfer network**, in which an edge means that agents may be moved directly between two tanks. The 20 tanks form 4 regions of 5 tanks. Each within-region pair is joined with probability $p_{in} = 0.6$ and each between-region pair with $p_{out} = 0.05$, so regions are internally dense and joined by a few bridge edges (Figure 1). A generated network is accepted only if it is connected, has at least one between-region edge, is not complete, and does not give every node the same betweenness; otherwise it is redrawn deterministically from the same seed. Network statistics are given in Appendix C.




Figure 1. Transfer network for network seed 200. Colour marks the region and node area is proportional to betweenness centrality. Dark edges join different regions. The black rings mark the two tanks that highest-betweenness quarantine closes (tanks 19 and 17).

Each day runs the same steps in a fixed order. First, quarantine starts or ends. Second, each agent attempts a transfer with probability `transfer_rate`. Attempts are processed in a random order and each is checked at once against the current tank states: a transfer is blocked if the origin is quarantined or if no neighbouring tank is both open and below capacity; otherwise the agent moves to an eligible neighbour chosen uniformly at random. Disease state does not affect movement. Third, mixing within a tank is complete, and a susceptible agent in tank j becomes infected with probability $1 - (1 - \text{beta})^{I_j}$, where I_j is the number of infectious agents in the tank after

movement. Fourth, each agent that was infectious at the start of the day recovers with probability γ and stays immune. Infections and recoveries are applied together, so a newly infected agent neither transmits nor recovers on the day it is infected. A run stops on the first day with no infectious agents, or is recorded as censored at day 365.

We compare three strategies: no intervention; **random quarantine**, which closes k tanks chosen uniformly at random; and **highest-betweenness quarantine**, which closes the k tanks with the highest normalised betweenness [8] in the pre-outbreak network (computed with NetworkX [12], which implements Brandes' algorithm [13]). Neither selector can see the infection. Both close $k = 2$ tanks (10% of the tanks, a small budget) on the same response day for $D = 14$ days, so each costs 28 tank-days, and the only difference between them is which tanks are chosen. A quarantined tank cannot send or receive transfers, but transmission and recovery inside it continue. The response delay is counted from the introduction of infection on day 0 and stands in for detection plus administrative response.

The disease parameters, the transfer-rate and delay levels and D were chosen in a two-stage pilot whose criteria were committed before it ran and which never compared the two strategies (Appendix B). It chose $\beta = 0.2$ and $\gamma = 0.1$ (mean infectious period 10 days) and transfer levels 0, 0.01, 0.025, 0.1, from no movement to a saturating epidemic. The delays, 1, 12, 33 days, are an immediate response, the median day infection first reaches a second tank, and the median time to peak. $D = 14$ days is the only candidate at least one infectious period long that still covers less than a quarter of the epidemic.

202 automated tests check invariants, extreme cases and a hand-traced three-tank scenario, and the invariants are also checked on every day of every run; no run failed (Appendix A).

3. Experimental design

The formal experiment crosses three factors: transfer rate (0, 0.01, 0.025, 0.1), response delay (1, 12, 33 days) and strategy. We used 20 network instances, each with 5 epidemic seeds of its own. Each combination of network, epidemic seed and transfer rate forms one **paired block**, giving 100 blocks per transfer rate. A block contains 13 runs: one no-intervention baseline shared by all delays and, for each delay, one betweenness run and 3 random runs with different policy seeds. This gives 5,200 runs in total.

Each run is driven by a network seed, an epidemic seed and, for random quarantine, a policy seed (Appendix G). Epidemic draws are *event-keyed*, an implementation of common random numbers [10]: the random number used for any (process, day, agent) event depends only on the epidemic seed and that key. Every strategy in a block therefore follows an identical trajectory until the quarantine starts.

The co-primary outcomes are the **final attack rate**, the share of agents ever infected, and the number of **affected tanks**, the tanks that ever held an infectious agent. For each block we take the difference *targeted* – *random*, averaging the random arm over its policy seeds (negative favours targeting), and report its mean and the relative reduction $-\text{mean}(\text{difference})/\text{mean}(\text{random})$, with 95% confidence intervals from a percentile bootstrap [14] that resamples whole networks, as recommended for clustered data [15] (2000 resamples). The comparison of each strategy with its block's baseline, the restriction to blocks in which quarantine started, and the nested-seed rerun of a first, crossed-seed run (Appendix D) were added after the formal results. Every number in this report is rendered from the result files (Appendix E).

4. Results

All 5,200 formal runs completed (failed: 0; censored: 0). In 0.0%, 2.0% and 12.0% of the non-zero-transfer blocks at delays of 1, 12 and 33 days, the outbreak died out before the quarantine started, and all strategies in those blocks are identical by construction.

4.1 Transfer rate drives the spread across regions

Without intervention, the transfer rate decides whether an outbreak stays in its first tank or spreads through the system (Table 1, Figure 2). With no movement, infection never leaves the initial tank: the mean attack rate is 0.049. The attack rate rises to 0.156 at a transfer rate of 0.01 (2.89 affected tanks) and to 0.519 at 0.025 (9.86 tanks); at 0.1 the outbreak saturates the system, with an attack rate of 0.972 and 19.44 of the 20 tanks affected.

Table 1. No-intervention outcomes by transfer rate (100 runs each; 95% CIs from a bootstrap that resamples whole networks).

Transfer rate	Final attack rate, mean [95% CI]	Affected tanks, mean [95% CI]	Peak infected, mean	Time to extinction, median (days)
0	0.049 [0.048, 0.050]	1.00 [1.00, 1.00]	8.59	27
0.01	0.156 [0.139, 0.174]	2.89 [2.55, 3.23]	14.94	49
0.025	0.519 [0.464, 0.567]	9.86 [8.78, 10.84]	29.71	90
0.1	0.972 [0.941, 0.999]	19.44 [18.86, 19.97]	85.61	73.5




Figure 2. Mean final attack rate (95% CI) by transfer rate and strategy, with one panel per response delay. The no-intervention baseline is the same in every panel. Affected tanks follow the same pattern (Appendix F, Figure F1).

Individual runs show how the spread happens (Figure 3). At transfer rate 0.025, infection never left its initial region in 21% of the baseline runs. In the representative cross-region outbreak, infection first peaks in its home region and reaches two further regions only weeks later.




Figure 3. Representative no-intervention runs at transfer rate 0.025: a local outbreak (network 218, epidemic seed 20093, attack rate 0.14) and a cross-region outbreak (network 211, epidemic seed 20059, attack rate 0.61). Each is the run whose attack rate is closest to the median of its class (local: 21 runs; cross-region: 79), a rule fixed before any trajectory was viewed. Top: S/I/R counts. Bottom: infectious agents per region.

4.2 Effect of quarantine and response delay

The quarantine reduces the attack rate only modestly, most at transfer rate 0.025 (Appendix F, Table F1). There, targeted quarantine cuts it by a similar amount at every delay (10.6%, 10.0% and 9.7% at 1, 12 and 33 days), whereas random quarantine fades (7.7%, 6.3% and 2.0%). At 0.01 and 0.1 the reductions are a few per cent or less. The intervals overlap across delays, so this contrast is descriptive.

4.3 Targeted versus random quarantine

Of the 18 paired comparisons with non-zero transfer (Table 2, Figure 4), 2 have a 95% CI that excludes 0, both at transfer rate 0.025 with a 33-day delay, where random quarantine had faded: the attack rate differs by -0.040 [-0.071, -0.009] (a relative reduction of 7.9%) and affected tanks by -0.83 [-1.46, -0.20]. Even there, targeting gives a lower attack rate in only 52% of blocks and a higher one in 33%. All other intervals include 0.

Table 2. Paired differences, targeted – random (negative favours targeted quarantine), with network-cluster bootstrap 95% CIs. The relative reduction is $-\text{mean}(\text{difference})/\text{mean}(\text{random})$. † marks a CI that includes 0.

Transfer rate	Delay (days)	Attack rate: targeted – random [95% CI]	Relative reduction	Affected tanks: targeted – random [95% CI]	Relative reduction
0.01	1	0.004 [-0.005, 0.014]†	-3.0%	0.10 [-0.08, 0.29]†	-3.8%
0.01	12	0.003 [-0.008, 0.016]†	-2.0%	0.09 [-0.12, 0.32]†	-3.1%
0.01	33	0.001 [-0.004, 0.007]†	-0.9%	0.04 [-0.06, 0.14]†	-1.3%
0.025	1	-0.015 [-0.058, 0.027]†	3.2%	-0.23 [-1.11, 0.65]†	2.5%
0.025	12	-0.019 [-0.058, 0.017]†	4.0%	-0.20 [-0.95, 0.52]†	2.2%
0.025	33	-0.040 [-0.071, -0.009]	7.9%	-0.83 [-1.46, -0.20]	8.6%
0.1	1	-0.006 [-0.036, 0.019]†	0.7%	-0.14 [-0.75, 0.37]†	0.7%
0.1	12	-0.003 [-0.020, 0.012]†	0.3%	-0.05 [-0.40, 0.25]†	0.2%
0.1	33	-0.005 [-0.012, 0.001]†	0.5%	-0.10 [-0.27, 0.03]†	0.5%




Figure 4. Mean paired difference, targeted – random, with 95% cluster-bootstrap CIs.

4.4 Why the quarantine effect is small (exploratory)

Why does a well-placed quarantine change so little? After the formal results, we measured three properties of the recorded runs (`scripts/mechanism_analysis.py`); this analysis is exploratory, not pre-registered.

Position. On the pre-outbreak networks, the two highest-betweenness tanks are endpoints of 40% of the between-region edges (of 8.1 per network), against 21% for the random pairs actually drawn, and removing them splits the network in 65% of networks, against 17%.

Timing. In the no-intervention runs at transfer rate 0.025, 21% of outbreaks never left their initial region; the others first reached a second region on day 15 (median; interquartile range 8–28). An immediate 14-day quarantine covers 49% of these first crossings, and 51% happen after it has ended; with a 12-day or 33-day delay, 39% or 80% have already happened when it starts. A fixed 14-day closure therefore overlaps only part of the period in which outbreaks cross regions.

Already-infected tanks. We had proposed that random quarantine fades because randomly chosen tanks are already infected by a late response. At a 33-day delay, 33% of random and 30% of targeted tanks had already held infection when quarantine started, so this explanation is not supported; the strategies differ in position, not in how often their tanks were already infected.

5. Discussion

H1 is supported. The no-intervention attack rate rises steeply with the transfer rate, and the steepest change falls between 0.01 and 0.025. Over this range, the share of outbreaks that reach a

second region rises from 32% to 79% (98% at 0.1), so outbreaks switch from mostly staying in one region to mostly crossing regions. This is consistent with the mobility threshold for invading new subpopulations in metapopulation models [3], although we did not estimate the threshold itself.

H2 is supported in direction, but the effect is small. The mean attack rate does not fall with a longer delay (apart from one dip under 0.01); at transfer rate 0.025 it rises under random quarantine from 0.479 to 0.508 but stays almost flat under targeted quarantine (0.464 to 0.468).

H3 receives limited support. Only one condition favours targeting with an interval excluding 0, and with 18 correlated comparisons about one in twenty would do so by chance. The data do not support a general advantage for targeting under this budget, only the narrower pattern that its effect did not fade with a later response.

Section 4.4 suggests why both strategies achieve so little. Targeting closes about twice as many bridges as a random choice; the constraint is time. Most outbreaks that spread first cross regions within about two weeks, so a 14-day closure either starts after the first crossing or ends before it. A possible reason, not tested here, why targeted quarantine kept its effect at a 33-day delay while random quarantine did not is that the bridge tanks can still delay the invasion of the regions that remain uninfected, whereas a random pair seldom lies on those paths. Salathé and Jones [5] found bridge-targeting effective for immunization, which is permanent; our results suggest that a temporary closure keeps this advantage only if it covers the period in which crossings occur. The analysis points to two testable improvements: a longer closure, and a closure triggered by the first detection outside the initial region rather than by a fixed delay.

These conclusions are limited in several ways. The turtles, tanks, disease and network are synthetic and uncalibrated, so magnitudes should not be read as predictions. We studied one quarantine budget and one disease regime, and the planned sensitivity analyses on β , γ and capacity were not run. There is no detection model. With 20 networks, the intervals are wide relative to the differences between strategies. An expanded 30-seed check of the network generator, done only after the formal experiment, found one exact tie for the second-highest betweenness, so the original no-tie acceptance criterion is not met (Appendix C). Finally, several choices were made after data were seen, including an added transfer level, a corrected pilot criterion for D , the nested-seed rerun, the comparison with baselines and the exploratory analysis of Section 4.4, whose timing measures come from the no-intervention runs; each is flagged in the text, recorded in Appendix B and the decision log, and the original results are kept.

6. Conclusion

In a modular tank system, the transfer rate decides whether a local outbreak stays local: raising it from 0.01 to 0.025 multiplies the attack rate by 3.3 and the number of affected tanks by 3.4. A quarantine of 2 tanks for 14 days reduces the attack rate by at most 10.6%, and targeting the bridge tanks is not consistently better than choosing at random, even though it closes about twice as many between-region edges. The exploratory analysis attributes this to timing: most crossings happen before or after a short, fixed-delay closure. Within this model, reducing cross-region transfers, or timing a closure to the first crossing, would matter more than which two tanks are closed; neither was tested here.

References

- [1] M. Gilbert, A. Mitchell, D. Bourn, J. Mawdsley, R. Clifton-Hadley and W. Wint, "Cattle movements and bovine tuberculosis in Great Britain," *Nature*, vol. 435, no. 7041, pp. 491–496, 2005. doi:10.1038/nature03548

- [2] R. Kao, L. Danon, D. Green and I. Kiss, "Demographic structure and pathogen dynamics on the network of livestock movements in Great Britain," *Proceedings of the Royal Society B*, vol. 273, no. 1597, pp. 1999–2007, 2006. doi:10.1098/rspb.2006.3505
- [3] V. Colizza and A. Vespignani, "Epidemic modeling in metapopulation systems with heterogeneous coupling pattern: Theory and simulations," *Journal of Theoretical Biology*, vol. 251, no. 3, pp. 450–467, 2008. doi:10.1016/j.jtbi.2007.11.028
- [4] M. J. Keeling and K. T. Eames, "Networks and epidemic models," *Journal of the Royal Society Interface*, vol. 2, no. 4, pp. 295–307, 2005. doi:10.1098/rsif.2005.0051
- [5] M. Salathé and J. H. Jones, "Dynamics and control of diseases in networks with community structure," *PLoS Computational Biology*, vol. 6, no. 4, e1000736, 2010. doi:10.1371/journal.pcbi.1000736
- [6] R. Pastor-Satorras and A. Vespignani, "Immunization of complex networks," *Physical Review E*, vol. 65, no. 3, 036104, 2002. doi:10.1103/PhysRevE.65.036104
- [7] R. M. Christley, G. L. Pinchbeck, R. G. Bowers, D. Clancy, N. P. French, R. Bennett and J. Turner, "Infection in social networks: Using network analysis to identify high-risk individuals," *American Journal of Epidemiology*, vol. 162, no. 10, pp. 1024–1031, 2005. doi:10.1093/aje/kwi308
- [8] L. C. Freeman, "A set of measures of centrality based on betweenness," *Sociometry*, vol. 40, no. 1, pp. 35–41, 1977. doi:10.2307/3033543
- [9] A. B. Downey, *Think Complexity: Complexity Science and Computational Modeling*, 2nd ed. Sebastopol, CA, USA: O'Reilly Media, 2018. ISBN 978-1-4920-4020-0
- [10] P. Glasserman and D. D. Yao, "Some guidelines and guarantees for common random numbers," *Management Science*, vol. 38, no. 6, pp. 884–908, 1992. doi:10.1287/mnsc.38.6.884
- [11] W. O. Kermack and A. G. McKendrick, "A contribution to the mathematical theory of epidemics," *Proceedings of the Royal Society of London. Series A*, vol. 115, no. 772, pp. 700–721, 1927. doi:10.1098/rspa.1927.0118
- [12] A. A. Hagberg, D. A. Schult and P. J. Swart, "Exploring network structure, dynamics, and function using NetworkX," in *Proceedings of the 7th Python in Science Conference*, pp. 11–15, 2008. doi:10.25080/TCWV9851
- [13] U. Brandes, "A faster algorithm for betweenness centrality," *Journal of Mathematical Sociology*, vol. 25, no. 2, pp. 163–177, 2001. doi:10.1080/0022250X.2001.9990249
- [14] B. Efron, "Bootstrap methods: Another look at the jackknife," *The Annals of Statistics*, vol. 7, no. 1, pp. 1–26, 1979. doi:10.1214/aos/1176344552
- [15] C. A. Field and A. H. Welsh, "Bootstrapping clustered data," *Journal of the Royal Statistical Society: Series B*, vol. 69, no. 3, pp. 369–390, 2007. doi:10.1111/j.1467-9868.2007.00593.x
- [16] C. R. Harris et al., "Array programming with NumPy," *Nature*, vol. 585, no. 7825, pp. 357–362, 2020. doi:10.1038/s41586-020-2649-2
- [17] W. McKinney, "Data structures for statistical computing in Python," in *Proceedings of the 9th Python in Science Conference*, pp. 56–61, 2010. doi:10.25080/Majora-92bf1922-00a
- [18] J. D. Hunter, "Matplotlib: A 2D graphics environment," *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007. doi:10.1109/MCSE.2007.55
- [19] M. E. J. Newman, "Modularity and community structure in networks," *Proceedings of the National Academy of Sciences*, vol. 103, no. 23, pp. 8577–8582, 2006. doi:10.1073/pnas.0601602103

Declaration: software, course material and AI tools

The model, experiments and analysis are implemented in Python 3.12 with NumPy [16], pandas [17], NetworkX [12] (whose betweenness function implements [13]) and Matplotlib [18]; exact versions are pinned in `requirements.txt`. The unit textbook [9] informed our understanding of networks and agent-based models, but none of its models or code is used; the tank, network and quarantine model, the experiments and the analysis were designed by the team for this project.

Generative AI tools were used in this project. Claude (Claude Code, Anthropic) assisted Member A with code, analysis scripts and the drafting and revision of documentation and of this report; Member B used ChatGPT (OpenAI) for parts of his code and documentation changes; and Codex CLI (OpenAI) performed read-only reviews of the code and results. The team reviewed all AI-assisted output, every number in this report is generated from the recorded results by the repository scripts, and every reference was checked against CrossRef, Semantic Scholar or the publisher record. The use of AI tools is also recorded in the repository README and contribution record.

Team contributions. Member A (Cam Zhou) wrote the research design, the baseline SIR model and event-keyed random draws, the network generator, the batch runner, the pilot and formal experiments, the analysis, the demonstration video and this report. Member B (Wenhao Zhang) wrote the run-record schema, the network-constrained movement and the two quarantine strategies, independently recalculated the hand-traced scenario, reviewed and corrected the pilot duration criterion, ran the expanded network audit, and wrote the project notebook. Both members reviewed each other's pull requests; the GitHub history and `docs/collaboration-plan.md` record who did what.

Appendix A. Verification

The test suite has 202 automated tests. They cover the invariants (agent count, $S + I + R = 200$, one tank per agent, capacity, no reinfection, quarantine blocking both directions, and reproducibility under identical seeds); extreme cases such as $\beta = 0$, $\gamma = 1$, $\text{transfer_rate} = 0$ and a full tank; the half-open quarantine interval and deterministic tie-breaking; and a 3-tank, 6-agent scenario traced by hand by both team members and matched against the model's event log. The invariants are also checked at runtime on every day of every run, and any violation ends the run as `failed`. No pilot or formal run failed.

Appendix B. Parameter selection (pilot)

β , γ , the transfer-rate levels, the response-delay levels and D were chosen in a two-stage pilot. Its selection criteria were committed to the repository before the pilot ran, and the pilot seeds were excluded from the formal experiment. The pilot never compared random with targeted quarantine.

Stage 1 used no-intervention runs only. It required, among other things, that with $\text{transfer_rate} = 0$ the infection never leaves the first tank; that the higher transfer levels do not consist almost entirely of minor outbreaks; that the lowest level does not saturate; and that three non-zero levels differ by at least a factor of two in accepted transfers per day.

In the first round (2,000 runs) none of the 8 (β , γ) candidates met the last criterion. The grid's adjacent levels differ by a nominal factor of exactly 2, and capacity blocking pulled the realised ratio below 2. As the protocol requires, we kept the threshold and reran the whole stage (2,400 runs) on an extended grid. The added level, 0.025, was chosen by one member after seeing the first round's transfer volumes; the protocol also requires both members to confirm a new grid, and this confirmation was still pending. 7 candidates then passed. The pre-registered selection rule picks the candidate whose median attack rate at the middle transfer level is closest to 0.5; it chose

$\beta = 0.2$ and $\gamma = 0.1$ (mean infectious period 10 days), with transfer levels 0, 0.01, 0.025, 0.1.

The three response delays are 1 (immediate), the median day on which infection first reaches a second tank, and the median time to peak. They were computed from 40 non-minor outbreaks and gave delays of 1, 12, 33 days. The pilot used a crossed seed design (10 epidemic seeds shared by 5 networks), so these quantities rest on only 10 independent outbreak origins.

Stage 2 considered $D = 7, 14$ and 21 days, with three acceptance criteria: (Q1) at least one attempted transfer is intercepted by a quarantine rule in $\geq 90\%$ of runs whose quarantine starts; (Q2) $D \leq 25\%$ of the median time to extinction; and (Q3) $D \geq$ the mean infectious period.

Under the original definition, which counted every blocked transfer in every run, no candidate passed Q1: 85.3% of runs had a blocked transfer. After inspecting the data we changed the denominator to runs whose quarantine started, giving 97.4%. Review then showed that this counter cannot measure Q1 at all. It also counts transfers blocked because every neighbouring tank was full, and in the no-intervention baselines 80.7% of runs already had such a block.

We therefore added three counters that attribute each blocked attempt to the first blocking rule reached: a quarantined origin; an open origin with no eligible destination but at least one quarantined neighbour with space; or remaining capacity/no-neighbour blocking. The definitions, the Q1 rule on them, and what would follow from each outcome were committed before the counters were implemented or the pilot rerun. The counters draw no random numbers, and the rerun reproduced all 27 pre-existing summary columns other than the run identifier, code commit and configuration hash across 5,550 runs exactly.

With these counters, the share of started runs with at least one attempt attributed to a quarantine rule was 91.8% / 95.7% / 96.7% for $D = 7 / 14 / 21$ days, so all candidates meet the operational Q1 criterion. For comparison, 90.6% of the same runs at $D = 14$ also had an attempt attributed to capacity. $D = 14$ is the only candidate that also satisfies Q2 ($D \leq 16.25$) and Q3, and the unchanged selection rule picks it. This is the duration the formal experiment had already used. The margin is narrowest at the lowest transfer level (0.01), where 90.9% of started runs had an attempt attributed to quarantine. The started-runs denominator remains a post-hoc choice.

The interpretation was clarified during the 7 October review without changing the counters or threshold: an attempted departure from a quarantined origin is counted even when all neighbours are full. Q1 measures interception by a rule, not whether removing quarantine would have made that particular attempt succeed, and not a causal reduction in infections. The Q1 check was measured only after the formal experiment had run, so a different result could not have changed the completed experiment.

Other post-hoc elements are the strategy-minus-baseline comparison, the rule for choosing representative runs, and the exploratory position and timing analysis of Section 4.4 (`scripts/mechanism_analysis.py`, written after the formal results to test the mechanism the earlier draft had only proposed). The second team member had not yet signed off the pilot protocol when the pilot ran, nor confirmed the extended transfer grid. All of these are recorded in `docs/decision-log.md`, and the original results are kept in the repository.

Appendix C. Network statistics

Across the 20 networks of the formal experiment:

Property	Median	Range
Modularity of the region partition [19]	0.51	0.31–0.61
Between-region edges	8	5–15
Mean degree	3.30	2.70–4.00
Diameter	6	4–9

3 networks needed at least one redraw.

The promised expanded structural check (network seeds 0-29) was completed only after the formal experiment. The selected setting meets the generation-stability criterion but fails the original no-rank-2-tie clause: at one seed, the two highest-betweenness tanks tie and are both selected when $k = 2$, so there is no tie across the selected/excluded boundary. The audit retains the failed criterion and the existing parameters; it does not alter the runtime tie rule or demonstrate robustness of the epidemic conclusions. Details are in `docs/network-audit-2026-10-08.md`.

Appendix D. Sensitivity to the seed design

A first formal run reused the same 5 epidemic seeds in every network, so the networks were crossed with the seeds instead of having seeds nested within them. Because draws are event-keyed, this gave every network the same initial infected agent and the same random draws, so outcomes were strongly correlated across networks; for example, whether an outbreak died out early depended on the seed rather than the network. The intended design nests epidemic seeds within networks. We found the problem after viewing that run's results, reran with nested seeds, and report the nested run. In the crossed run, 7 of the 18 targeted – random comparisons had a CI excluding 0, against 2 in the nested design: with only five shared outbreak origins, the crossed run overstated its precision.

The local/cross-region classification of outbreaks uses end-of-day snapshots, so an infectious agent that arrives in a tank and recovers on the same day is not seen. This happened in 5 of 400 baseline runs. The 2 of them classed as local were rerun, and their affected tanks were all inside the initial region.

Appendix E. Reproduction

All results can be regenerated from the repository:

```
python scripts/run_experiment.py experiments/config/formal-nested.json
python scripts/analyse_results.py formal-nested
python scripts/build_report.py
```

Appendix F. Additional results

Table F1. Relative reduction in mean final attack rate compared with the block's own no-intervention baseline (descriptive; added after the formal results). † marks a mean-difference CI that includes 0.

Transfer rate	Delay (days)	Random: reduction in attack rate [95% CI]	Betweenness: reduction in attack rate [95% CI]
0.01	1	5.3% [-0.6%, 10.5%]†	2.4% [-5.2%, 9.4%]†
0.01	12	3.7% [-0.6%, 8.8%]†	1.8% [-3.9%, 7.6%]†
0.01	33	2.1% [0.4%, 4.5%]	1.3% [-1.6%, 4.4%]†
0.025	1	7.7% [-0.6%, 14.6%]†	10.6% [2.5%, 18.6%]
0.025	12	6.3% [-1.1%, 13.2%]†	10.0% [0.6%, 19.5%]
0.025	33	2.0% [-3.5%, 7.6%]†	9.7% [3.6%, 16.4%]
0.1	1	1.5% [0.6%, 2.6%]	2.2% [-0.0%, 5.2%]†
0.1	12	1.7% [0.7%, 2.9%]	2.0% [0.5%, 3.9%]
0.1	33	0.3% [-0.1%, 0.7%]†	0.8% [0.1%, 1.6%]




Figure F1. Mean number of affected tanks, laid out as in Figure 2.

Appendix G. Parameters and seeds

Table G1. Parameters of the formal experiment. "Pilot" values were selected by the pre-registered pilot (Appendix B); "fixed" values were set at Checkpoint 1 and recorded in `docs/decision-log.md`.

Parameter	Value	Meaning	Source
Agents / tanks / regions	200 / 20 / 4	population and housing structure	fixed
Tank capacity	12	maximum agents per tank	fixed (D002)
Initial occupancy	10 per tank, 1 infected	start state; the infected agent is drawn from the epidemic seed	fixed
p_in/p_out	0.6 / 0.05	edge probability within / between regions	structural audit (D005)
beta	0.2	daily infection probability per infectious tank-mate	pilot
gamma	0.1	daily recovery probability (mean infectious period 10 days)	pilot
transfer_rate	0, 0.01, 0.025, 0.1	daily probability that an agent attempts a transfer	pilot
Response delay	1, 12, 33 days	days from introduction to the start of quarantine	pilot
k / D	2 tanks / 14 days	quarantine budget (28 tank-days)	fixed (D003) / pilot (D004)
Horizon	365 days	runs still infectious at this day are censored	fixed (D006)
Network seeds	200-219	one network per seed; redrawn up to 100 times if a structural check fails	design
Epidemic seeds	20000-20099	5 per network, nested	design
Policy seeds	1000-1002	random quarantine choices	design
Bootstrap	2000 resamples	network-cluster percentile CIs	analysis